

Lab04

In this lab, yolov4 was implemented using DarkNet as its backbone. Since the given config file as well as the architecture of the DarkNet were specifically for yolov3, a new config file and some modifications to the DarkNet architecture were necessary. The config file for yolov4 and the pretrained weights were taken from <https://github.com/AlexeyAB/darknet/blob/master/cfg/yolov4-csp.cfg> and https://github.com/AlexeyAB/darknet/releases/download/darknet_yolo_v3_optimal/yolov4.weights respectively.

Apart from the those, some modifications of the DarkNet architecture are as follows:

Maxpool (Ref: Prof. Matt)

(in forward:)

```
if module_type == "convolutional" or module_type == "upsample" or module_type == "maxpool":
```

```
    x = self.module_list[i]
```

(in create_modules:)

```
    elif x["type"] == "maxpool":
```

```
        stride = int(x["stride"])
```

```
        size = int(x["size"])
```

```
        assert size % 2
```

```
        maxpool = nn.MaxPool2d(kernel_size=size, stride=stride, padding=size // 2)
```

```
        module.addmodule("maxpool{0}".format(index), maxpool)
```

Mish activation

```
class Mish(nn.Module):
```

```
    def init(self):
```

```
        super().init()
```

```
    def forward(self, x):
```

```
        return x * torch.tanh(F.softplus(x))
```

In create module

```
    elif activation == "mish":
```

```
        activn = Mish()
```

```
        module.addmodule("mish{0}".format(index), activn)
```

The input image size from 416 x 416 to 608 x 608

The route (Ref: Prof. Matt)

```
    elif module_type == "route":
```

```
        layers = module["layers"]
```

```
        layers = [int(a) for a in layers]
```

```
    if (layers[0]) > 0:
```

```
        layers[0] = layers[0] - i
```

```

if len(layers) == 1: # 1 item in layer

    x = outputs[i + (layers[0])]
else:# more than 1 item in layer
    if len(layers) == 4:# 4 items in layer

        if (layers[1]) > 0:
            layers[1] = layers[1] - i
        if (layers[2]) > 0:
            layers[2] = layers[2] - i
        if (layers[3]) > 0:
            layers[3] = layers[3] - i
        map1 = outputs[i + layers[0]]
        map2 = outputs[i + layers[1]]
        map3 = outputs[i + layers[2]]
        map4 = outputs[i + layers[3]]
        x = torch.cat((map1, map2, map3, map4), 1)
    else: # 2 items in layer

        if (layers[1]) > 0:
            layers[1] = layers[1] - i
        map1 = outputs[i + layers[0]]
        map2 = outputs[i + layers[1]]
        x = torch.cat((map1, map2), 1)

```

Conversion of colors in prep_images

```

def prep_image(img, inp_dim):
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = letterbox_image(img, (inp_dim, inp_dim))
    return torch.from_numpy(img.transpose(2, 0, 1)).float().div(255.0).unsqueeze(0)

```

As for the inference, I have tested yolov4 with some COCO images, the results can be seen below.

Note that I have also converted the images from RGB back to BGR since during the prep_image the images were converted from BGR to RGB.

```

In [2]: from __future__ import division
import time
import torch
import torch.nn as nn
from torch.autograd import Variable
import numpy as np
import cv2
from util import *
import argparse
import os
import os.path as osp
from darknet import Darknet
import pickle as pkl
import pandas as pd

```

```

import random

images = "cocoimages"
batch_size = 4
confidence = 0.5
nms_thesh = 0.4
start = 0
CUDA = torch.cuda.is_available()

num_classes = 80
classes = load_classes("data/coco.names")

#Set up the neural network

print("Loading network.....",file=open("output_coco.txt", "a"))
model = Darknet("cfg/yolov4.cfg")
model.module_list[114].conv_114 = nn.Conv2d(2048, 512, kernel_size=(1, 1), st
model.load_weights("yolov4.weights")
print("Network successfully loaded",file=open("output_coco.txt", "a"))

model.net_info["height"] = 416
inp_dim = int(model.net_info["height"])
assert inp_dim % 32 == 0
assert inp_dim > 32

#If there's a GPU availible, put the model on GPU

if CUDA:
    model.cuda()

# Set the model in evaluation mode

model.eval()

read_dir = time.time()

# Detection phase

try:
    imlist = [osp.join(osp.realpath('.'), images, img) for img in os.listdir(
except NotADirectoryError:
    imlist = []
    imlist.append(osp.join(osp.realpath('.'), images))
except FileNotFoundError:
    print ("No file or directory with the name {}".format(images),file=open("
    exit()

if not os.path.exists("des"):
    os.makedirs("des")

load_batch = time.time()
loaded_ims = [cv2.imread(x) for x in imlist]

im_batches = list(map(prepare_image, loaded_ims, [inp_dim for x in range(len(iml
im_dim_list = [(x.shape[1], x.shape[0]) for x in loaded_ims]
im_dim_list = torch.FloatTensor(im_dim_list).repeat(1,2)

leftover = 0
if (len(im_dim_list) % batch_size):
    leftover = 1

if batch_size != 1:
    num_batches = len(imlist) // batch_size + leftover
    im_batches = [torch.cat((im_batches[i*batch_size : min((i + 1)*batch_size,

```

```

len(im_batches))])) for i in range(num_batches)]

write = 0

if CUDA:
    im_dim_list = im_dim_list.cuda()

start_det_loop = time.time()
for i, batch in enumerate(im_batches):
    # Load the image
    start = time.time()
    if CUDA:
        batch = batch.cuda()
    with torch.no_grad():
        prediction = model(Variable(batch), CUDA)

    prediction = write_results(prediction, confidence, num_classes, nms_conf

    end = time.time()

    if type(prediction) == int:

        for im_num, image in enumerate(imlist[i*batch_size: min((i + 1)*batch_size, len(imlist))]):
            im_id = i*batch_size + im_num
            print("{0:20s} predicted in {1:6.3f} seconds".format(image.split("/")[-1], end=""), file=open("output_coco.txt", "a"))
            print("{0:20s} {1:s}".format("Objects Detected:", " "), file=open("output_coco.txt", "a"))
            print("-----")
            continue

    prediction[:,0] += i*batch_size #transform the attribute from index in imlist to index in im_batches

    if not write: #If we have't initialised output
        output = prediction
        write = 1
    else:
        output = torch.cat((output,prediction))

    for im_num, image in enumerate(imlist[i*batch_size: min((i + 1)*batch_size, len(imlist))]):
        im_id = i*batch_size + im_num
        objs = [classes[int(x[-1])] for x in output if int(x[0]) == im_id]
        print("{0:20s} predicted in {1:6.3f} seconds".format(image.split("/")[-1], end=""), file=open("output_coco.txt", "a"))
        print("{0:20s} {1:s}".format("Objects Detected:", " ").join(objs), file=open("output_coco.txt", "a"))
        print("-----", file=open("output_coco.txt", "a"))

    if CUDA:
        torch.cuda.synchronize()

try:
    output
except NameError:
    print ("No detections were made",file=open("output_coco.txt", "a"))
    exit()

im_dim_list = torch.index_select(im_dim_list, 0, output[:,0].long())

scaling_factor = torch.min(416/im_dim_list,1)[0].view(-1,1)

output[:,[1,3]] -= (inp_dim - scaling_factor*im_dim_list[:,0].view(-1,1))/2
output[:,[2,4]] -= (inp_dim - scaling_factor*im_dim_list[:,1].view(-1,1))/2

output[:,1:5] /= scaling_factor

for i in range(output.shape[0]):
    output[i, [1,3]] = torch.clamp(output[i, [1,3]], 0.0, im_dim_list[i,0])
    output[i, [2,4]] = torch.clamp(output[i, [2,4]], 0.0, im_dim_list[i,1])

```

```

output_recast = time.time()
class_load = time.time()
colors = [[255, 0, 0], [255, 0, 0], [255, 255, 0], [0, 255, 0], [0, 255, 255]]

draw = time.time()

def write(x, results):
    c1 = tuple(x[1:3].int())
    c2 = tuple(x[3:5].int())
    img = results[int(x[0])]
    cls = int(x[-1])
    color = random.choice(colors)
    label = "{0}".format(classes[cls])
    cv2.rectangle(img, c1, c2, color, 1)
    t_size = cv2.getTextSize(label, cv2.FONT_HERSHEY_PLAIN, 1, 1)[0]
    c2 = c1[0] + t_size[0] + 3, c1[1] + t_size[1] + 4
    cv2.rectangle(img, c1, c2, color, -1)
    cv2.putText(img, label, (c1[0], c1[1] + t_size[1] + 4), cv2.FONT_HERSHEY_PLAIN, 1, color)
    return img

list(map(lambda x: write(x, loaded_ims), output))

det_names = pd.Series(imlist).apply(lambda x: "{}/det_{}".format("des", x.split('.')[0]))

list(map(cv2.imwrite, det_names, loaded_ims))

end = time.time()

print("SUMMARY", file=open("output_coco.txt", "a"))
print("-----", file=open("output_coco.txt", "a"))
print("{:25s}: {}".format("Task", "Time Taken (in seconds)"), file=open("output_coco.txt", "a"))
print()
print("{:25s}: {:.23f}".format("Reading addresses", load_batch - read_dir), file=open("output_coco.txt", "a"))
print("{:25s}: {:.23f}".format("Loading batch", start_det_loop - load_batch), file=open("output_coco.txt", "a"))
print("{:25s}: {:.23f}".format("Detection (" + str(len(imlist)) + " images)", end - load_batch), file=open("output_coco.txt", "a"))
print("{:25s}: {:.23f}".format("Output Processing", class_load - output_recast), file=open("output_coco.txt", "a"))
print("{:25s}: {:.23f}".format("Drawing Boxes", end - draw), file=open("output_coco.txt", "a"))
print("{:25s}: {:.23f}".format("Average time_per_img", (end - load_batch)/len(imlist)), file=open("output_coco.txt", "a"))
print("-----", file=open("output_coco.txt", "a"))

torch.cuda.empty_cache()


```

In [18]:

```

import matplotlib.pyplot as plt
import matplotlib.image as mpimg
image = mpimg.imread("des/det_dog-cycle-car.png")
plt.imshow(image)
plt.show()

```



Inference on COCO dataset

In [1]:

```
results_COCO = open("output_coco.txt", "r")
print(results_COCO.read())
```

Loading network.....

Network successfully loaded

dog-cycle-car.png predicted in 0.015 seconds

Objects Detected: bicycle truck dog

000000016502.jpg predicted in 0.015 seconds

Objects Detected: sheep

000000016598.jpg predicted in 0.015 seconds

Objects Detected: person tie cell phone

000000016958.jpg predicted in 0.015 seconds

Objects Detected: chair chair chair vase vase

000000017029.jpg predicted in 0.015 seconds

Objects Detected: car car dog frisbee

000000017031.jpg predicted in 0.015 seconds

Objects Detected: person giraffe

000000017115.jpg predicted in 0.015 seconds

Objects Detected: zebra zebra

000000017178.jpg predicted in 0.015 seconds

Objects Detected: car horse horse horse

000000017182.jpg predicted in 0.015 seconds

Objects Detected: apple chair book book book

000000017207.jpg predicted in 0.015 seconds

Objects Detected: person car motorbike bus truck

000000017379.jpg predicted in 0.015 seconds

Objects Detected: person tvmonitor sink sink

000000017436.jpg predicted in 0.015 seconds

Objects Detected: person bench

000000017714.jpg predicted in 0.016 seconds

```

Objects Detected:      cup cup fork fork knife bowl banana banana diningtable
-----
0000000017899.jpg      predicted in  0.016 seconds
Objects Detected:      person cup cup bowl sandwich sandwich chair chair sofa di
ningtable
-----
0000000017905.jpg      predicted in  0.016 seconds
Objects Detected:      person traffic light
-----
0000000017959.jpg      predicted in  0.016 seconds
Objects Detected:      person person kite kite kite kite
-----
0000000018150.jpg      predicted in  0.009 seconds
Objects Detected:      person person bottle pizza
-----
0000000018193.jpg      predicted in  0.009 seconds
Objects Detected:      person donut chair
-----
SUMMARY
-----
Task                    : Time Taken (in seconds)
Reading addresses       : 0.001
Loading batch           : 0.123
Detection (18 images)   : 0.478
Output Processing       : 0.000
Drawing Boxes           : 0.142
Average time_per_img    : 0.041

```

In [8]:

```

import cv2
import matplotlib.pyplot as plt
import glob
for img in glob.glob("des/*.jpg"):
    cv_img = cv2.imread(img)
    cv_img = cv2.cvtColor(cv_img, cv2.COLOR_RGB2BGR)
    plt.imshow(cv_img)
    plt.show()

```

